

THE PITFALLS OF DEEP LEARNING IN REAL-WORLD APPLICATIONS: A CASE STUDY

Anonymous authors

Paper under double-blind review

ABSTRACT

Deep learning has shown remarkable success in various domains, yet its deployment in real-world applications often encounters significant challenges. This paper investigates the pitfalls of deep learning models when applied to practical scenarios, focusing on issues such as data availability, model interpretability, and computational requirements. Through a series of experiments, we demonstrate that even state-of-the-art models can fail to generalize effectively, leading to sub-optimal performance. Our findings highlight the importance of addressing these challenges to ensure the reliable deployment of deep learning in real-world applications.

1 INTRODUCTION

Deep learning has revolutionized many fields, from computer vision to natural language processing Goodfellow et al. (2016). However, its application in real-world scenarios often reveals significant limitations. These limitations include difficulties in handling sparse or noisy data, challenges in model interpretability, and high computational costs. This paper explores these pitfalls through a series of experiments designed to simulate real-world conditions. Our results show that even well-established models can struggle to generalize effectively, underscoring the need for robust solutions to these challenges.

2 RELATED WORK

Previous research has highlighted the challenges of deploying deep learning in real-world applications. provides a comprehensive analysis of these challenges, focusing on issues such as data availability and model interpretability. compares classical machine learning and deep learning models, revealing limitations in deep learning approaches for anomaly detection. Additionally, discusses the importance of interpretability and reliability in deep learning models. Our work builds on these studies by empirically demonstrating the pitfalls of deep learning in practical scenarios.

3 BACKGROUND

Deep learning models, particularly those based on neural networks, have achieved state-of-the-art performance in various tasks. However, their success often relies on large amounts of high-quality data and substantial computational resources Goodfellow et al. (2016). In real-world applications, these conditions are rarely met, leading to challenges such as overfitting, poor generalization, and difficulty in model interpretation. Understanding these limitations is crucial for developing more robust and reliable deep learning systems.

4 METHOD

We conducted a series of experiments to evaluate the performance of deep learning models under real-world conditions. We used a standard neural network architecture trained with the Adam optimizer. The models were evaluated on datasets with varying levels of noise and sparsity to simulate

real-world challenges. Our evaluation metrics included loss and accuracy, which were measured across multiple epochs to assess the models' generalization capabilities.

5 EXPERIMENTAL SETUP

Our experimental setup involved training neural networks on datasets designed to mimic real-world conditions. We used a standard training procedure with the Adam optimizer and evaluated the models' performance on separate validation sets. The datasets were intentionally corrupted with varying levels of noise and sparsity to simulate the challenges encountered in practical applications. We measured the models' loss and accuracy over multiple epochs to assess their ability to generalize.

6 EXPERIMENTS

Our experiments revealed significant challenges in the deployment of deep learning models in real-world scenarios. As shown in Figure 1, the models struggled to achieve low loss values, particularly when trained on noisy data. Figure 2 further illustrates the impact of data quality on model performance, with higher average loss observed for models trained on corrupted datasets. Additionally, Figure 3 shows that the models' accuracy plateaued early, indicating limited generalization capability.




Figure 1: Loss vs Epoch Comparison Across Methods




Figure 2: Comparison of Average Loss Across Methods




Figure 3: Accuracy vs Epoch Comparison Across Methods

7 CONCLUSION

Our experiments highlight the significant challenges of deploying deep learning models in real-world applications. Despite their success in controlled environments, these models often struggle to generalize effectively when faced with noisy or sparse data. Our findings underscore the need for robust solutions to these challenges, including improved data preprocessing techniques and more interpretable model architectures. Future work should focus on addressing these limitations to ensure the reliable deployment of deep learning in practical scenarios.

REFERENCES

Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT Press, 2016.

SUPPLEMENTARY MATERIAL

A APPENDIX SECTION

In this appendix, we provide additional details on our experimental setup and results. This includes hyperparameter settings, dataset descriptions, and additional figures illustrating the models’ performance on different datasets.