

THE PITFALLS OF DEEP LEARNING: A CASE STUDY ON OPTIMIZATION FAILURES AND DATASET BIASES

Anonymous authors

Paper under double-blind review

ABSTRACT

Despite the widespread success of deep learning across various domains, significant challenges remain in ensuring robust and reliable model performance. This paper explores two critical pitfalls: optimization failures and dataset biases. Through a series of experiments, we demonstrate how these issues can lead to suboptimal model performance and discuss potential mitigation strategies. Our findings highlight the importance of addressing these challenges to improve the generalizability and fairness of deep learning models in real-world applications.

1 INTRODUCTION

Deep learning has revolutionized many fields, from computer vision to natural language processing (Goodfellow et al., 2016). However, deploying these models in real-world scenarios often reveals significant challenges, including optimization failures and dataset biases (?). These issues can lead to models that perform poorly on unseen data or exhibit unfair behavior towards certain groups. In this paper, we present a case study that examines these pitfalls in detail, providing insights into their impact and potential mitigation strategies.

2 RELATED WORK

Previous research has extensively documented the challenges of deep learning models in real-world applications. ? reviewed the limitations of deep learning in complex traffic environments, emphasizing issues such as generalizability and explainability. ? highlighted vulnerabilities to adversarial attacks, while ? proposed defenses against such attacks. Additionally, ? explored the impact of demographic biases in datasets and proposed fairness-aware mitigation techniques. Our work builds on these studies by focusing specifically on optimization failures and dataset biases, providing empirical evidence of their impact.

3 BACKGROUND

Optimization failures, such as vanishing gradients and overfitting, are well-documented challenges in training deep learning models (?). These issues can prevent models from learning effectively, leading to poor performance. Dataset biases, on the other hand, can result in models that unfairly favor certain groups or fail to generalize to new data (?). Understanding and addressing these challenges is crucial for developing robust and fair deep learning models.

4 METHOD

We conducted a series of experiments to explore the impact of optimization failures and dataset biases on model performance. Our methodology involved training deep learning models on both synthetic and real-world datasets, while systematically varying optimization parameters and dataset composition. We evaluated model performance using standard metrics, such as accuracy and fairness, and analyzed the results to identify patterns and trends.

5 EXPERIMENTAL SETUP

Our experiments were conducted using a variety of deep learning architectures, including convolutional neural networks (CNNs) and recurrent neural networks (RNNs). We used both synthetic datasets, designed to simulate specific biases, and real-world datasets, such as ImageNet. Optimization parameters, including learning rate and batch size, were varied to induce different types of optimization failures. We also manipulated the composition of the training datasets to introduce biases, such as demographic imbalances.

6 EXPERIMENTS

Our experiments revealed significant challenges associated with optimization failures and dataset biases. Figure 1 shows the impact of varying optimization parameters on model performance, highlighting the sensitivity of deep learning models to these settings. Figure 2 compares the performance of models trained on biased and unbiased datasets, demonstrating the detrimental effect of dataset biases on fairness and generalizability.




Figure 1: Impact of optimization parameters on model performance.



Figure 2: Comparison of model performance on biased and unbiased datasets.

7 CONCLUSION

Our findings underscore the importance of addressing optimization failures and dataset biases in deep learning. These challenges can significantly impact model performance and fairness, limiting the effectiveness of deep learning in real-world applications. Future work should focus on developing robust optimization techniques and fairness-aware data processing methods to mitigate these issues. By addressing these pitfalls, we can improve the reliability and fairness of deep learning models, enabling their successful deployment in diverse and challenging environments.

REFERENCES

Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT Press, 2016.

SUPPLEMENTARY MATERIAL

A APPENDIX SECTION

Additional details on the experimental setup, including hyperparameters and dataset descriptions, are provided in this appendix. We also include supplementary figures and tables that further illustrate our findings.